


# FINAL REPORT

## RAFT CONSENSUS ALGORITHM

---



Joachim Kirkegaard Friis s093256  
Anders Nielsen s103457

November 14, 2014

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Problem . . . . .	3
1.2	Motivation . . . . .	4
1.2.1	The Two Generals Problem . . . . .	4
1.3	Raft . . . . .	5
<b>2</b>	<b>Requirements</b>	<b>6</b>
<b>3</b>	<b>Analysis</b>	<b>7</b>
3.1	Impossibility of the Two Generals Problem . . . . .	7
3.2	Raft as a solution . . . . .	7
<b>4</b>	<b>Design</b>	<b>8</b>
<b>5</b>	<b>Implementation</b>	<b>9</b>
<b>6</b>	<b>Tests</b>	<b>10</b>
6.1	Deterministic vs. in deterministic tests . . . . .	10
<b>7</b>	<b>Conclusion</b>	<b>11</b>

# 1 Introduction

This is the final report for a project done in the course "Fault tolerant systems" 02228. The purpose of this report is to document an implementation of a consensus algorithm i.e. Raft.

Starting, the fundamental problem when talking about consensus in a distributed system, will be presented. This will then present the motivation behind the project itself. As many solutions to this problem already exist, it should also be discussed why Raft in particular is relevant in the context of this project.

## 1.1 Problem

Reaching consensus in a distributed system means that processes in the network agree on some value of state of the entire system. This is often needed as processes might be faulty and thus reliability and availability is at stake on single-point-of-failure. Upholding these properties in a given distributed system then relies on the architecture utilised and hardware implemented in the processes. And having a system that can tolerate faulty processes by maintaining a common knowledge of a value of state is preferred. A simple solution to this could be to initiate a vote among all correct processes on to what the value is, in which the value is the result of the majority vote. The figure 1 below illustrates an example of a distributed system consisting of a number of nodes. The value  $x$  is what the system must agree upon. Though here we have a faulty process  $P1$ , taking the fact of connectivity of system aside, the system will suffer from this because of the unreliable messages transmitted from this process might override this value at some of the other correct ones.

The fundamental problem behind this, is that you cannot rely on the individual process

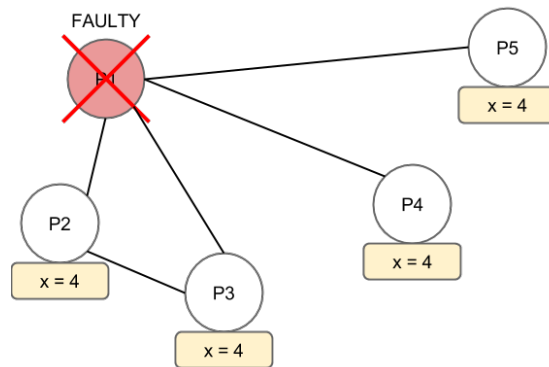


Figure 1: A distributed system consisting of a number of nodes with a faulty one.

to be reliable and thus solely store the value. This means, that every process should store this value and should be able to be altered somehow. This boils down to a well-known problem - The Two Generals Problem.

## 1.2 Motivation

The elements of the problem presented by this problem serves as motivation behind consensus in a system of unreliable components. Because how do you know for certain what a value is, when you for sure know that some components must be faulty at some point.

### 1.2.1 The Two Generals Problem

Agreeing on which state a system, among its processes, is in is a well known problem when talking about distributed systems. This can be illustrated by the Two Generals Problem analogy. Here we see two generals of the same army who want to attack an

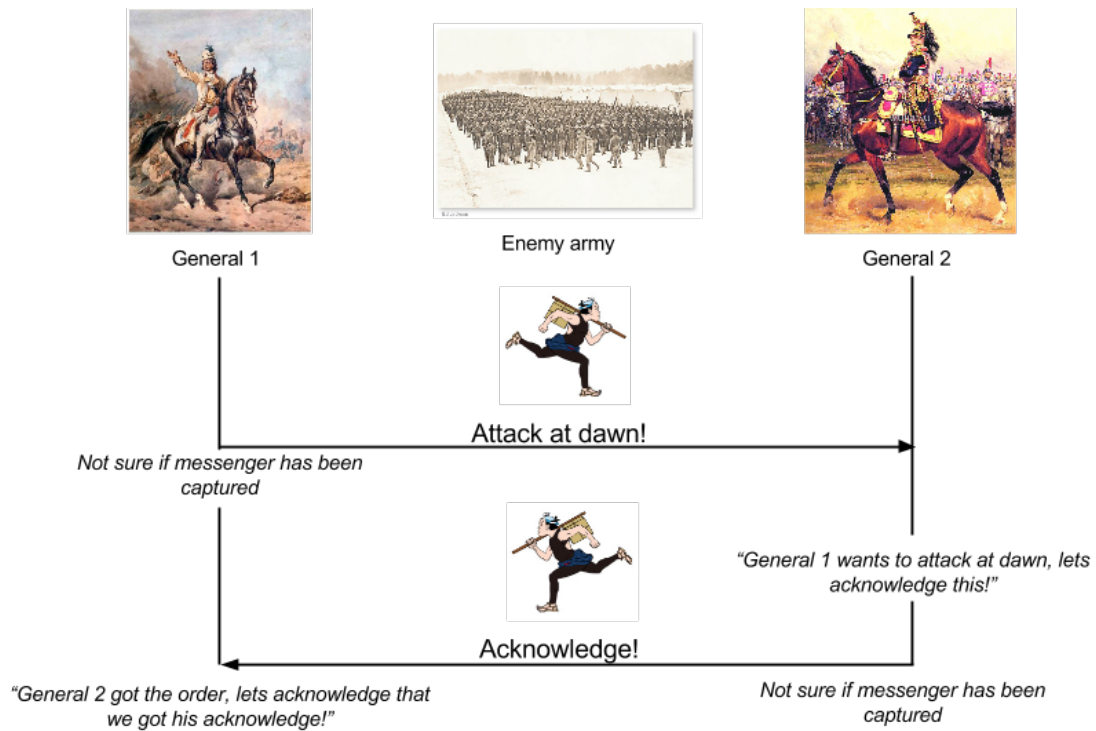


Figure 2: Two generals tries to agree on when to attack the enemy by sending a messenger, but they are not sure the messenger survives his trip between their camps.

enemy army. But they have to attack at the same time in order to win. They cant communicate directly to each other since they are at different fronts of the battlefield i.e. in the own camps. In the context of a distributed system, the generals here can be seen as two processes trying to agree on a value. General 1 then sends out a messenger to tell the second general, that they should attack at dawn. The second general then receives this message, but the first general cannot be sure of this (the messenger might be captured or killed by the enemy on his way to the second general and vice versa).

Again, in the context of distributed system, the unreliability of messenger can directly related back to the unreliability of message transmission in a normal distributed system. So the second general sends the messenger back in order to acknowledge this. But the first general also has to acknowledge this, resulting in a never ending run for the poor messenger - thus the generals can never agree on when to attack the enemy.

### **1.3 Raft**

## 2 Requirements

The tool we are implementing should serve as an illustration tool, to show how the Raft algorithm works for given scenarios of faulty processes in a distributed system.

- a. The tool must illustrate a scenario of a given set of processes in a distributed system, in which a leader is elected.
- b. The user must then be able to disconnect the leader.
  - a) If a leader is disconnected a new leader must be elected automatically.
- c. The user must be able to add new processes to the system.
- d. The user should be able to input parameters to vary the scenario, such as the number of processes, and the heart rate of the system.
  - a) The system should be able to have the parameters changed at run-time.
- e. The tool must be an implementation of the Raft algorithm.
  - a) The Raft implementation should have the following properties: Safety, availability, timing independence, and that commands complete as soon as a majority has responded.
- f. The tool could be further extended with a visualisation of the given distributed system.

### **3 Analysis**

*In this section we will come back to the Two Generals problem and further analyse and realize that reaching consensus without a leader/authority, who can issue commands, is impossible. Then we will present Raft as a solution to this problem and compare it to other solutions and discuss why we have chosen it.*

#### **3.1 Impossibility of the Two Generals Problem**

#### **3.2 Raft as a solution**

## 4 Design

*In this section we will relate to our requirements of the tool, and from this design the implementation.*



## 5 Implementation

## 6 Tests

### 6.1 Deterministic vs. in deterministic tests

- a. Simulation and validation of different scenarios described in the requirements.
  - a) Leader election:
    - i. The system must choose a new leader if the current leader is disconnected.
    - ii. The system must be able to choose a leader if the system is partitioned.
  - b) Timing and availability
    - i. The system must not depend on timing of new elections, thus not producing incorrect results because of unknown events.

## 7 Conclusion

## References

- [1] Michael Bailey, Evan Cooke, Farnam Jahanian, Yunjing Xu, and Manish Karir. A survey of botnet technology and defenses. In *Proceedings of the 2009 Cybersecurity Applications & Technology Conference for Homeland Security*, CATCH '09, pages 299–304, Washington, DC, USA, 2009. IEEE Computer Society.
- [2] John Black, Martin Cochran, and Trevor Highland. A study of the md5 attacks: Insights and improvements. In Matthew Robshaw, editor, *Fast Software Encryption*, volume 4047 of *Lecture Notes in Computer Science*, pages 262–277. Springer Berlin Heidelberg, 2006.
- [3] David Dagon, Guofei Gu, and ChristopherP. Lee. A taxonomy of botnet structures. In Wenke Lee, Cliff Wang, and David Dagon, editors, *Botnet Detection*, volume 36 of *Advances in Information Security*, pages 143–164. Springer US, 2008.
- [4] William Rex. The zombies all around you. <http://www.meannnasty.com/zombies-around/>, april 2014.